Moderations

Explore how to use OpenAI's moderations endpoint to identify and flag harmful or non-compliant text content. Understand request and response parameters, moderation models, and practical use cases such as filtering hate speech, violence, and self-harm topics in real-world applications.

The moderations endpoint

Moderations is an endpoint that OpenAI provides to flag text. This API endpoint analyzes the text provided and evaluates whether the content complies with Open AI's usage policies. After the API analyzes the text, it tells the user which flags fit the text and the confidence that the API has for each flag.

Moderations

Ask

Moderations API call

To analyze a chunk of text, we can call the following function:

```
response = client.moderations.create(input="Text to analyze.")
```

Understanding the moderations endpoint

Let’s look at the moderation endpoint in more detail, reviewing the various models, the request parameters, and the response parameters.

Moderation models

Moderation contains two models that achieve the same result. However, they differ in their approach to updates. Let's examine the two models.

Model	Description
`text-moderation-latest`	This model is used to flag text. It is automatically updated o updates arrive.
`text-moderation-stable`	This model is used to flag text as well. It is not updated auto notification is sent to the user.

Request parameters

Let’s look at the parameters that are required to make a request at the moderations endpoint.

Fields	Format	Type	Description
`model`	String	Optional	The ID of the engine to u the request.
`input`	String or array	Required	The text that would be fl the pre-determined cate the API.

Ask

> **Note**: You can learn more about the `model` argument in this lesson.

Response parameters

Let's look at the response given by the Moderation API and evaluate the information each parameter provides.

Object	Format	Description
`id`	String	The unique identifier of the response to your request.
`model`	String	The model that was used to generate the results.
`results`	Array	The results contain three objects: • `flagged` (Boolean): This is a simple classification of v violates Open AI's usage policy. • `categories` (Object): These are all the pre-determine text is evaluated against. If any of these are violated, t `true` . • `category_scores` (Object): These are the confidence determined categories that the text is evaluated agair

Let's take a look at all the categories that the API endpoint can flag.

> **Note:** Any given statement can have one or many flags be marked as `true`.

Category	Description
Hate	This category is marked `true` when the statement expresses hate for any group based on their gender, race, ethnicity, reli orientation, disability status or caste.

Ask

Hate/threatening	This category is marked `true` when the statement contains to violence towards a group based on their gender, race, eth nationality, sexual orientation, disability status or caste.
Harassment	This category is marked `true` when the statement uses, inc harassing language towards anyone.
Harassment/threatening	This category is marked `true` when the statement contains anyone.
Self-harm	This category is marked `true` when the statement promote shows any act of self-harm.
Self-harm/intent	This category is marked `true` when the statement has the u or engagement in self-harm acts.
Self-harm/instructions	This category is marked `true` when the statement provides encourages acts of self-harm.
Sexual	This category is marked `true` when the statement tries to c sexual activity or promotes sexual services (this does not inc and wellness).
Sexual/minors	This category is marked `true` when the statement contains under 18 years old.
Violence	This category is marked `true` when the statement depicts c physical injury.
Violence, graphic	This category is marked `true` when the statement depicts c physical injury in a graphic manner.

The input to this endpoint is a text string that contains the request. The output is also a list that contains all the categories mentioned above marked as `true` or `false`.

Ask

Moderations: input and output model

Flagging through moderation

Using the moderations endpoint, developers can transmit a single flag to indicate that specific content may potentially violate OpenAI's usage policies. This flag serves as a concise notification, enabling efficient identification and assessment of the content in question.

By utilizing the moderation endpoint to transmit this single flag, developers can prompt appropriate action, such as implementing filters or conducting further review, to ensure compliance with OpenAI's guidelines and standards. This streamlined process enhances content moderation efforts and facilitates the maintenance of a safe and compliant platform environment.

Let's explore some of the use cases where we transmit a single flag using the moderations endpoint.

Hate speech

Using the moderations endpoint, we have the capability to identify hate speech within a given prompt. Our approach involves supplying an input to the mo subsequently analyzing the outcomes to ascertain the presence of hate speech.

Ask

Please provide values for the following:

SECRET_KEY Not Specified...

Edit Import Values from JSON

main.py Python 3.8

```
from openai import OpenAI
client = OpenAI(api_key="{{Secret_key}}")

response = client.moderations.create(input="I hate them")

output = response.results[0]
print(output)
```

Run

Hate speech detection

Violence

The moderations endpoint can be used to identify content that exhibits violent themes, including depictions of death, physical harm, or acts of violence.

Please provide values for the following:

SECRET_KEY Not Specified...

Edit Import Values from JSON

Ask

main.py Python 3.8

```
from openai import OpenAI
client = OpenAI(api_key="{{Secret_key}}")

response = client.moderations.create(input="The unidentified joker will b

output = response.results[0]
print(output)
```

Violence detection

Self-harm

The moderations endpoint can be used to detect content that promotes, encourages, or depicts acts of self-harm, such as suicide, cutting, and eating disorders.

Please provide values for the following:

SECRET_KEY Not Specified...

Edit Import Values from JSON

main.py Python 3.8

```
from openai import OpenAI
client = OpenAI(api_key="{{SECRET_KEY}}")

response = client.moderations.create(input="Joker hurts himself when he l

output = response.results[0]
```

Ask

```
print(output)
```

Run

Self harm detection

Practical use case

OpenAI's moderations endpoint can be employed for a variety of real-world applications across different industries and domains. Some examples include:

- **Social media platforms:** Social media companies can utilize the moderations endpoint to automatically detect and filter out content containing violence, hate speech, or graphic imagery, thereby fostering a safer and more inclusive online environment for their users.
- **E-commerce websites:** Online marketplaces and retailers can use the moderations endpoint to scan product listings and user reviews for fraudulent or misleading content, as well as to identify and remove listings for prohibited items such as weapons or counterfeit goods.
- **News and media outlets:** News organizations and publishers can employ the moderations endpoint to moderate user comments on articles and posts, filtering out spam, hate speech, or abusive content to maintain a respectful and constructive discourse among readers.

Ask